

Gaurav Kumar

Data Analyst / Python Developer / ML Engineer

Bengaluru, India | +91 7481812280 | gauravbaba311@gmail.com

linkedin.com/in/gauravkumar311 | github.com/GauravBaba3 | gaurav-kumar-311.netlify.app

SUMMARY

Computer Science graduate with proven ability to build LLM-powered applications, end-to-end machine learning pipelines, and interactive BI dashboards. Delivered two AI-driven tools using NLP and Large Language Models as a fresher, alongside production-deployed Flask web apps. Proficient in Python, SQL, Power BI, Scikit-learn, and Django. Actively seeking entry-level Data Analyst or Python Developer roles to deliver measurable, data-driven impact.

SKILLS

Programming Languages: Python, SQL, JavaScript

Data & ML Libraries: Pandas, NumPy, Matplotlib, Seaborn, Scikit-learn

Machine Learning: Supervised Learning, Random Forest, Classification, Feature Engineering, Hyperparameter Tuning, Cross-Validation, EDA, ETL Pipelines, Data Modeling, Data Wrangling

AI / NLP: Large Language Models (LLM), NLP Techniques, TF-IDF Vectorization, Cosine Similarity, Generative AI Fundamentals, Hugging Face

Web Frameworks: Django, Streamlit, Bootstrap, Tailwind CSS, HTML, CSS, REST APIs, JSON

Data Visualization & BI: Power BI, MS Excel (Advanced), Power Query, KPI Dashboards, Data Storytelling

Databases: MySQL, Oracle SQL, Query Optimization

Tools & Platforms: Git, GitHub, Jupyter Notebook, VS Code, AWS IoT Analytics

WORK EXPERIENCE

Data Science and Full-Stack Web Development Trainee

Aug 2025 – Present

Qspiders, Bengaluru

- Performed EDA on 10,000+ structured records across sales and HR domains using Python and SQL, surfacing 5+ actionable business insights per dataset cycle.
- Designed and automated KPI reporting dashboards in Power BI, cutting manual reporting effort by 40% and enabling real-time performance tracking across 3 business units.
- Optimized SQL queries with joins, aggregations, and window functions, reducing average query execution time by 20% on multi-table datasets.
- Built Python-based data cleaning pipelines to resolve missing values, duplicates, and outliers, improving downstream model input quality by standardizing 98%+ of records.
- Created stakeholder-facing visualizations using Matplotlib and Seaborn that communicated trend analysis and forecasts to non-technical team leads.

Web Development Intern

Feb 2025 – Mar 2025

NSIC Technical Services Centre, Chennai

- Developed 5+ responsive web pages using HTML, CSS, JavaScript, and Bootstrap, improving page load times by 30% and ensuring cross-device compatibility.
- Integrated Bootstrap components into the internal portal UI, reducing user friction by 25% based on internal usability testing across the team.
- Built dynamic features including form validations and navigation bars, enhancing usability for 200+ monthly portal users.
- Implemented MySQL database queries for backend data handling, contributing to a 15% improvement in server response time through query restructuring.
- Resolved 20+ front-end bugs during structured code reviews, delivering production-ready releases across 2 sprint cycles.

PROJECTS

AI Resume Screening & Top-N Candidate Ranking System

GitHub

- Built an end-to-end NLP pipeline that extracts skills, experience, and keywords from PDF resumes and ranks candidates against a job description using TF-IDF vectorization and SentenceTransformer, achieving **88% matching accuracy** on a 150-resume test set.
- Engineered a Scikit-learn Pipeline with ColumnTransformer and hyperparameter tuning (GridSearchCV) to optimize the ranking model; reduced false-positive shortlisting by 22%.
- Deployed an interactive Streamlit web interface integrating an LLM for contextual relevance scoring, enabling HR teams to screen 10x faster than manual review.
- Technologies:* Python, Streamlit, LLM, TF-IDF, SentenceTransformer, Scikit-learn, Pandas, PyPDF2, Pipeline, GridSearchCV

Natural Language Query to Data Filtering & Visualization System

GitHub

- Developed a Streamlit application that translates plain-English queries (e.g., “show top 5 products by revenue in Q3”) into structured Pandas filters using an LLM, tested on a 50,000-row retail dataset.
- Generated dynamic bar, line, and pie charts from filtered results, reducing analysis time for non-technical users by an estimated 60% vs. manual filtering.

- Integrated LLM query parsing with structured JSON output mapping to Pandas operations, achieving **91% query interpretation accuracy** across 200+ test queries.
- *Technologies: Python, Streamlit, LLM, Pandas, Matplotlib, JSON, Encoding & Preprocessing Pipelines*

Breast Cancer Prediction System

GitHub

- Trained a Random Forest classifier on the Breast Cancer Dataset (569 samples, 30 features) with 5-fold cross-validation, achieving **96.5% accuracy, and 0.96 F1-score** on the held-out test set.
- Performed feature selection using correlation analysis and SelectKBest, reducing dimensionality from 30 to 18 features while maintaining model performance.
- Deployed the model as a Flask REST API with an HTML/CSS/Bootstrap front-end, enabling real-time predictions via a user-facing web form.
- *Technologies: Python, Scikit-learn, Pandas, NumPy, Flask, REST API, HTML, CSS, Bootstrap*

Blinkit Sales Analytics Dashboard

GitHub

- Designed a Power BI dashboard tracking **14 KPIs** across product categories, outlet types, and sales regions, enabling data-driven stock and pricing decisions.
- Wrote optimized SQL queries with CTEs and window functions for data extraction, reducing dashboard refresh time by 35%.
- *Technologies: Power BI, SQL, DAX, Data Modeling, KPI Reporting*

EDUCATION

Bachelor of Engineering (B.E.) – Computer Science and Engineering

2021 – 2025

Anna University, Tamil Nadu

CGPA: 7.57 / 10

- Coursework in Data Structures, Algorithms, Database Management, Machine Learning, and Software Engineering.
- Final-year project: Road Crash Injury Severity Prediction using ML models.

Higher Secondary (12th Standard)

2019 – 2021

Mahesh Singh Yadav College, Gaya

63%

CERTIFICATIONS

- **Data Analytics Job Simulation** – Deloitte (Forage): Applied analytics to solve structured business problems including data interpretation and insight delivery.
- **Data Analytics & Visualisation Simulation** – Accenture (Forage): End-to-end data cleaning, visualisation, and stakeholder storytelling across multiple datasets.
- **Software Engineering Job Simulation** – JPMorgan Chase & Co. (Forage): REST API development and system integration using Kafka and H2 database.
- **Fundamentals of Generative AI** – Microsoft Learn: Core LLM concepts, AI-generated content, and real-world generative AI applications.
- **AutoML for Production Model Development** – Microsoft Learn: Automated model selection, training, and deployment using Azure AutoML.